

Team: Shock Waves

Members: Nilanjana Chakraborty 23BAI1073

C. Varshah 23BAI1371

PID controller code:

```
import gym
import pybullet as p
import numpy as np
import time
import csv

from gym_pybullet_drones.envs.HoverAviary import HoverAviary

env = HoverAviary(gui=True)
env.reset()

def generate_random_waypoints(num_waypoints, x_range, y_range, z_range):
    waypoints = []
    for _ in range(num_waypoints):
        x = np.random.uniform(*x_range)
        y = np.random.uniform(*y_range)
        z = np.random.uniform(*z_range)
        waypoints.append([x, y, z])
    return waypoints

num_waypoints = 5
x_range = (-2, 2)
y_range = (-2, 2)
z_range = (1, 3)

waypoints = generate_random_waypoints(num_waypoints, x_range, y_range, z_range)

with open('pid_values.csv', mode='w', newline='') as file:
    writer = csv.writer(file)

    writer.writerow(['Waypoint', 'Step', 'Error_X', 'Error_Y', 'Error_Z', 'Action_X', 'Action_Y', 'Action_Z',
                    'Action_W'])

    def navigate_to_waypoints(env, waypoints, max_steps=500):
```

```

for waypoint in waypoints:

    target_x, target_y, target_z = waypoint

    step_counter = 0

    print(f" Moving to waypoint: {waypoint}")

    while step_counter < max_steps:

        step_counter += 1

        action = np.zeros((1, 4))

        obs, reward, done, info, _ = env.step(action)

        print(f" Observation received: {obs} (Shape: {obs.shape})")

        if obs.size >= 3:

            pos = np.array(obs[0, :3], dtype=float)

        else:

            print(f"Invalid observation size: {obs.size}, assigning default position.")

            pos = np.array([0, 0, 1.5])

        error = np.array([target_x - pos[0], target_y - pos[1], target_z - pos[2], 0])

        random_scale = np.random.uniform(0.1, 0.3)

        base_action = np.clip(error * random_scale + 0.5, 0.4, 1.6)

        noise = np.random.normal(0, 0.05, base_action.shape) # Add some noise

        action = base_action + noise

        action = np.reshape(action, (1, 4))

        print(f" Action Shape: {action.shape} | Action: {action}")

        writer.writerow([waypoint, step_counter, error[0], error[1], error[2], action[0, 0], action[0,
1], action[0, 2], action[0, 3]])

        env.step(action)

        env.render()

        if np.linalg.norm(error[:3]) < 0.05: # Only check position, not yaw

            print(f" Waypoint {waypoint} reached!")

            break

        time.sleep(0.1)

navigate_to_waypoints(env, waypoints)

```

```
env.close()
```

RL model code:

```
import torch

import numpy as np

from stable_baselines3 import PPO, DDPG, SAC

from gym_pybullet_drones.envs.HoverAviary import HoverAviary

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

print(f"Using device: {device}")

env = HoverAviary(gui=False)

env.reset()

TIMESTEPS = 200000 # Increase for better learning

MODEL_SAVE_PATH = "hoveraviary_model"

models = {

    "PPO": PPO(

        "MlpPolicy", env, verbose=1, device=device,

        n_steps=8192, batch_size=2048, learning_rate=0.0005, gamma=0.995,

        ent_coef=0.01, vf_coef=0.5, max_grad_norm=0.5

    ),

    "DDPG": DDPG(

        "MlpPolicy", env, verbose=1, device=device,

        learning_rate=0.001, batch_size=1024, gamma=0.99, tau=0.005

    ),

    "SAC": SAC(

        "MlpPolicy", env, verbose=1, device=device,

        learning_rate=0.0003, batch_size=512, gamma=0.98, tau=0.005, ent_coef="auto"

    ),

}

for algo, model in models.items():

    print(f" Training {algo} model...")

    model.learn(total_timesteps=TIMESTEPS)
```

```

model.save(f"{MODEL_SAVE_PATH}_{algo}")

print(f"{algo} model saved successfully!\n")

for algo in models.keys():

    loaded_model = models[algo].load(f"{MODEL_SAVE_PATH}_{algo}", env)

    print(f"{algo} model loaded successfully!")

print("All models trained, saved, and ready to use!")

```

Training process screenshots:

rollout/	
ep_len_mean	15.8
ep_rew_mean	18.7
time/	
fps	43
iterations	2
time_elapsed	376
total_timesteps	16384
train/	
approx_kl	0.010764032
clip_fraction	0.0565
clip_range	0.2
entropy_loss	-5.68
explained_variance	-0.0146
learning_rate	0.0005
loss	35.6
n_updates	10
policy_gradient_loss	-0.0134
std	1
value_loss	73.4

rollout/	
ep_len_mean	16.9
ep_rew_mean	19.8
time/	
fps	47
iterations	3
time_elapsed	517
total_timesteps	24576
train/	
approx_kl	0.012233392
clip_fraction	0.0614
clip_range	0.2
entropy_loss	-5.68
explained_variance	-0.00831
learning_rate	0.0005
loss	36.5
n_updates	20
policy_gradient_loss	-0.0131
std	1
value_loss	77.6

rollout/	
ep_len_mean	21.6
ep_rew_mean	23.7
time/	
fps	37
iterations	4
time_elapsed	865
total_timesteps	32768
train/	
approx_kl	0.012279859
clip_fraction	0.0662
clip_range	0.2
entropy_loss	-5.66
explained_variance	0.0161
learning_rate	0.0005
loss	37.6
n_updates	30
policy_gradient_loss	-0.0143
std	0.994
value_loss	86.8

CSV File:

1	Waypoint	Step	Error_X	Error_Y	Error_Z	Action_X	Action_Y	Action_Z	Action_W
2	[-0.08933853983005413, 0.9168321518824785, 1.2075776892672374]	1	-0.08933854	0.916832152	1.095077692	0.4491114	0.768551292	0.842139328	0.53722549
3	[-0.08933853983005413, 0.9168321518824785, 1.2075776892672374]	2	-0.089396216	0.916847077	1.093955337	0.57863088	0.612602174	0.706943451	0.435093525
4	[-0.08933853983005413, 0.9168321518824785, 1.2075776892672374]	3	-0.090096748	0.917016686	1.091522387	0.571055653	0.77462541	0.755332209	0.438396884
5	[-0.08933853983005413, 0.9168321518824785, 1.2075776892672374]	4	-0.092600999	0.917504215	1.087759181	0.521155067	0.648863921	0.769526586	0.483493834
6	[-0.08933853983005413, 0.9168321518824785, 1.2075776892672374]	5	-0.098653358	0.918348845	1.082805118	0.558553446	0.637346223	0.609775718	0.51217742
7	[-0.08933853983005413, 0.9168321518824785, 1.2075776892672374]	6	-0.11055353	0.919605241	1.076974778	0.515461732	0.564120179	0.539965779	0.493795509
8	[-0.08933853983005413, 0.9168321518824785, 1.2075776892672374]	7	-0.130883842	0.921336491	1.070802121	0.417106669	0.506839054	0.615569023	0.485612105
9	[-0.08933853983005413, 0.9168321518824785, 1.2075776892672374]	8	-0.16233944	0.923591881	1.064975544	0.361824383	0.676950258	0.679214926	0.503485882
10	[-0.08933853983005413, 0.9168321518824785, 1.2075776892672374]	9	-0.207812502	0.926600738	1.06030093	0.429983488	0.710207297	0.684419716	0.577812383
11	[-0.08933853983005413, 0.9168321518824785, 1.2075776892672374]	10	-0.270567461	0.930854126	1.057958587	0.473282111	0.743745848	0.801202541	0.376718838
12	[-0.08933853983005413, 0.9168321518824785, 1.2075776892672374]	11	-0.354173361	0.937046073	1.059737458	0.416674646	0.635893542	0.692571485	0.632440839
13	[-0.08933853983005413, 0.9168321518824785, 1.2075776892672374]	12	-0.462526022	0.945958038	1.06819184	0.401303214	0.718311139	0.663890612	0.552066346
14	[-0.08933853983005413, 0.9168321518824785, 1.2075776892672374]	13	-0.599498509	0.958590469	1.086652762	0.453482113	0.609394921	0.638694799	0.502910851
15	[-0.08933853983005413, 0.9168321518824785, 1.2075776892672374]	14	-0.768640338	0.976096544	1.119274652	0.371831648	0.627791692	0.634850038	0.452569523
16	[-0.08933853983005413, 0.9168321518824785, 1.2075776892672374]	15	-0.954835175	0.999872427	1.13422238	0.440351132	0.548331538	0.713126871	0.592208091
17	[-0.08933853983005413, 0.9168321518824785, 1.2075776892672374]	16	-1.136023282	1.027196537	1.142858415	0.426439822	0.782178184	0.790738971	0.539530198
18	[-0.08933853983005413, 0.9168321518824785, 1.2075776892672374]	17	-1.276204346	1.051087263	1.13492678	0.398921911	0.805381595	0.789015684	0.42904001
19	[-0.08933853983005413, 0.9168321518824785, 1.2075776892672374]	18	-1.428832768	1.080428469	1.142933516	0.392946082	0.668312624	0.687165029	0.551198785
20	[-0.08933853983005413, 0.9168321518824785, 1.2075776892672374]	19	-1.57027769	1.106987181	1.181505619	0.453122178	0.6744945	0.713198954	0.484560915

For the PID controller, we generated 70 rows of data.